



PromCon

North America 2021

pktvisor.dev

Deep network traffic observability with pktvisor and Prometheus

Shannon Weyrick • VP Research • Office of CTO

sweyrick@ns1.com

The Big Picture



PromCon
North America 2021

pktvisor and **Orb**
supplement modern
observability stacks
by facilitating
**edge network
observability**

The projects are **free**
and **open source**
backed by **NSI**.


The ultimate
goal is **dynamic
orchestration** of
business intelligence
at the **edge**

1. Deep Network Observability

2.  pktvisor

3.  pktvisor +
Prometheus

4.  Orb



Deep Network Observability

Deep Network Observability



PromCon
North America 2021



- 🌀 Unwrapping and inspecting network traffic and activity can provide insight useful for **operations**, **debugging** and **security**
- 🌀 But **analysis** and **collection** of traffic across a distributed and often ephemeral set of end points is **hard**
- 🌀 How do we orchestrate and **extract insights** from these flows?

Context From NS1.



PromCon
North America 2021

- 🌀 At **NS1.** we run Managed DNS and other critical network services across many globe spanning networks
- 🌀 We need to **tune** our global anycast networks for the **best delivery time**
- 🌀 We are often subject to **malicious traffic** which we need to understand to be able to **protect against**
- 🌀 We need to **debug individual delivery** nodes at high resolution
- 🌀 We need a **global view** of all nodes and to drill in to **different dimensions of network traffic data** over time

What Do We Want To Know?



PromCon
North America 2021

- 🔗 What are the **counters**, **rates** and **frequent items** across common network traffic dimensions?
- 🔗 How many **unique** IP addresses and query names (**cardinality**) are there?
- 🔗 What are the important **quantiles** of transaction timings? What's the **histogram** of response payloads?
- 🔗 What is the **amplification factor** from Query to Response size?
- 🔗 What is **still querying** that DNS record that was **deleted**?
- 🔗 From what **ASN** and **Geo** regions is traffic coming from?
- 🔗 Is this traffic spike **malicious** or **legitimate**? Is this a random label attack? Is it widely distributed? IPv4? UDP? Against what zone?



pktvisor

“packet visor”

The Big Picture: pktvisor




PromCon
North America 2021

- 🌀 **Free** and **open source** observability **agent** backed by NSI
- 🌀 Cut its teeth **observing critical infrastructure** for the past 7 years
- 🌀 Currently packet capture focused, more **input sources** on tap
- 🌀 Goals
 - Support **pluggable input sources** and **analyzers** (contributors!)
 - Drive observability via **dynamic policies** over **REST API**
 - Support **modern** observability stacks

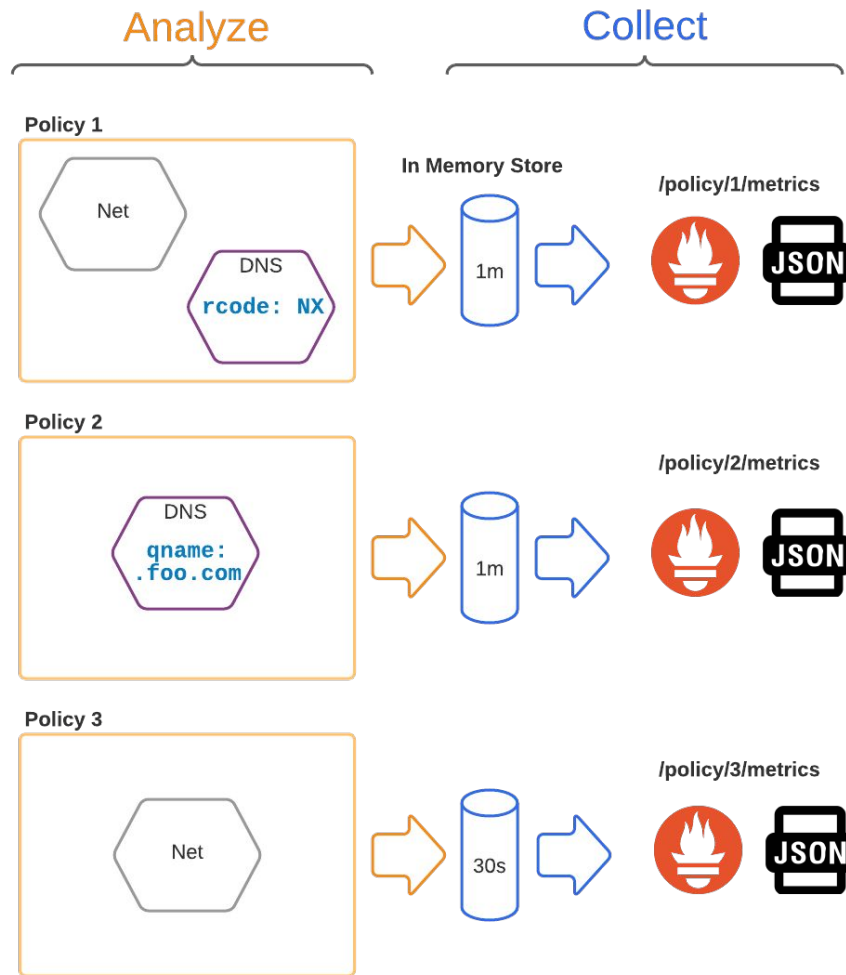

pktvisor.

← 32 bits →										
ver	hlen	TOS		pkt len						
identification				flg	fragment offset					
TTL		protocol		header cksum						
Source IP address										
Destination IP address										
Source port				Destination port						
UDP length				UDP cksum						
Query ID		q	r	opcode	a	a	r	r	z	rcode
Question count				Answer count						
Authority count				Addl. Record count						
DNS question or answer data										

DNS packet on the wire



pktvisor.



```

sweyrick@blesk:~$ curl -w "\n" -s localhost:10853/api/v1/policies/default/metrics/bucket/1 | jq . -C | head -39
{
  "default-default_dns": {
    "dns": {
      "cardinality": {
        "qname": 40
      },
      "period": {
        "length": 60,
        "start_ts": 1633445201
      },
      "rates": {
        "total": {
          "p50": 0,
          "p90": 14,
          "p95": 24,
          "p99": 52
        }
      },
      "top_nxdomain": [
        {
          "estimate": 1,
          "name": "ns-1022.awsdns-63.ne"
        },
        {
          "estimate": 1,
          "name": "lb._dns-sd._udp.0.1.168.192.in-addr.arpa"
        }
      ],
      "top_qname2": [
        {
          "estimate": 20,
          "name": ".malwarebytes.com"
        },
        {
          "estimate": 20,
          "name": ".apple.com"
        },
        {
          "estimate": 18,

```

```

sweyrick@blesk:~$ curl -w "\n" -s localhost:10853/api/v1/policies/default/metrics/prometheus | highlight -O xterm256 -S bash | head -22
# HELP packets_rates_pps_in Rate of ingress in packets per second
# TYPE packets_rates_pps_in summary
packets_rates_pps_in{instance="blesk",module="default-default_net",policy="default",quantile="0.5"} 4
packets_rates_pps_in{instance="blesk",module="default-default_net",policy="default",quantile="0.9"} 13
packets_rates_pps_in{instance="blesk",module="default-default_net",policy="default",quantile="0.95"} 14
packets_rates_pps_in{instance="blesk",module="default-default_net",policy="default",quantile="0.99"} 26
packets_rates_pps_in_sum{instance="blesk",module="default-default_net",policy="default"} 26
packets_rates_pps_in_count{instance="blesk",module="default-default_net",policy="default"} 59
# HELP packets_rates_pps_out Rate of egress in packets per second
# TYPE packets_rates_pps_out summary
packets_rates_pps_out{instance="blesk",module="default-default_net",policy="default",quantile="0.5"} 3
packets_rates_pps_out{instance="blesk",module="default-default_net",policy="default",quantile="0.9"} 16
packets_rates_pps_out{instance="blesk",module="default-default_net",policy="default",quantile="0.95"} 22
packets_rates_pps_out{instance="blesk",module="default-default_net",policy="default",quantile="0.99"} 31
packets_rates_pps_out_sum{instance="blesk",module="default-default_net",policy="default"} 31
packets_rates_pps_out_count{instance="blesk",module="default-default_net",policy="default"} 59
# HELP packets_rates_pps_total Rate of all packets (combined ingress and egress) in packets per second
# TYPE packets_rates_pps_total summary
packets_rates_pps_total{instance="blesk",module="default-default_net",policy="default",quantile="0.5"} 6
packets_rates_pps_total{instance="blesk",module="default-default_net",policy="default",quantile="0.9"} 29
packets_rates_pps_total{instance="blesk",module="default-default_net",policy="default",quantile="0.95"} 34
packets_rates_pps_total{instance="blesk",module="default-default_net",policy="default",quantile="0.99"} 57
sweyrick@blesk:~$ curl -s localhost:10853/api/v1/policies/default/metrics/prometheus | highlight -O xterm256 -S bash | head -22

```

pktvisor-cli (client: 3.3.0-develop | server: 3.3.0-develop)

Pkts 3249 | UDP 1126 (34.7%) | TCP 2073 (63.8%) | Other 50 (1.5%) | IPv4 3196 (98.4%) | IPv6 3 (0.1%) | In 1629 (50.9%) | Out 1570 (49.1%) | Deep Samples 3249 (100.0%)

Pkt Rates Total 5/s 6/20/24/114 pps | In 3/s 3/9/12/59 pps | Out 2/s 3/10/14/61 pps | IP Card. In: 272 | Out: 284 | TCP Errors 0 | OS Drops 0 | IF Drops 0

DNS Wire Pkts 1112 (34.2%) | Rates Total 0/s 0/0/0/0 | UDP 1112 (100.0%) | TCP 0 (0.0%) | IPv4 1112 (100.0%) | IPv6 0 (0.0%) | Query 566 (50.9%) | Response 546 (49.1%)

DNS Xacts 546 | Timed Out 20 | In 176 (32.2%) | Out 370 (67.8%) | In 30.5/1627.4/1815.9/3220.9 ms | Out 21.7/82.4/94.1/260.3 ms | Qname Card. 252

DNS NOERROR 542 (99.3%) | SRVFAIL 0 (0.0%) | NXDOMAIN 4 (0.7%) | REFUSED 0 (0.0%) | Time Window 5:43PM to 5:48PM, Period 297s

Top QName 2 .amazon.com 118 (10.6%) .cloudfront.net 68 (6.1%) .roku.com 56 (5.0%) .co.uk 48 .google.com 40 .amazonaws.com 39 .awsstatic.com 22	Top QName 3 .logs.roku.com 56 (5.0%) .us-east-1.amazonaws.com 27 (2.4%) .dscg.akamaiedge.net 16 (1.4%) www.shopbop.com 13 .shortbread.aws.dev 12 .amazon-blogs.psdops.com 12 www.imdb.com 10	Top NX lb._dns-sd._udp.0.1.168.192.in-addr.arpa 2 rns-779.awsdns-33.net 1 (0.2%) us-east-1.console.aws.amazon.com 1 (0.2%)	Slow In p3epsfumangb5dvijf52er7wdi.appsync-api.us- www.imdb.com 3 (0.5%) d23tl967axkois.cloudfront.net 2 (0.4%) player.live-video.net 2 sparrow.wondershare.com 2 d2in0p32vp1pij.cloudfront.net 2 blog.aboutamazon.com 2
Top QTypes A 1060 (95.3%) HTTPS 38 (3.4%) AAAA 10 (0.9%) PTR 4	Top RCodes NOERROR 542 (99.3%) NXDOMAIN 4 (0.7%)	Top SRVFAILS	Slow Out ns10.tmobileus.net 3 (0.5%) gtm-cn-v64163wlk09.gtm-a2b4.com 2 (0.4%) cache.prod.amazon-blogs.psdops.com 2 (0.4%) prod.log.shortbread.aws.dev 2 www.shopbop.com 2 www.zappos.com 2 filmstock-api-eus.wondershare.cc 2
Top REFUSED	IPv4 192.168.0.217 2132 (65.6%) 34.102.140.197 242 (7.4%) 34.122.121.32 35 (1.1%) 192.43.172.30 20 192.42.93.30 18 172.253.122.188 16 208.78.70.31 14	IPv6 ff02::1:2 3 (0.1%)	Top DNS UDP Ports 10139 4 (0.4%) 7096 4 (0.4%) 28397 4 (0.4%) 13766 4 10702 2 37127 2 14127 2
Top GeoLoc Unknown 2146 (66.1%) NA/United States 931 (28.7%) NA/United States/CA/Mountain View 32 (1.0%) EU 14 NA/United States/CA/San Jose 9 EU/Ireland/L/Dublin 8 AS/China 6	Top ASN Unknown 2150 (66.2%) 16509/AMAZON-02 361 (11.1%) 15169/GOOGLE 325 (10.0%) 397213/ULTRADNS 47 10515/CLT-NIC 36 397215/ULTRADNS 35 33517/DYNDNS 32		

Command Line UI

Easy Docker Install

ns1labs/pktvisor



PromCon
North America 2021

pull the image

```
root@host:~$ docker pull ns1labs/pktvisor
```

start the agent with a default collection policy

```
root@host:~$ docker run --net=host -d ns1labs/pktvisor pktvisord --prometheus eth0
```

scrape prometheus metrics

```
root@host:~$ curl localhost:10853/metrics
```

run the command line UI (ctrl-c to quit)

```
root@host:~$ docker run -it --rm --net=host ns1labs/pktvisor pktvisor-cli
```

pktvisor



<https://pktvisor.dev>

- FOSS
- Contributors welcome
- Please star!

ns1labs / pktvisor Public

Unwatch 42

Star 357

Fork 25

<> Code Issues 28 Pull requests 1 Discussions Actions Projects 2 Wiki Security Insights Settings

develop 5 branches 10 tags

Go to file Add file Code

weyrick feature/cli refactor (#120) cd0fed1 10 days ago 281 commits

.github/workflows	no automatic build for master	5 months ago
3rd	switch random number generators used by sampling (#110)	3 months ago
RFCs	require a new key 'kind' to specify policy type, in preparation for f...	13 days ago
appimage	require binary arg to appimage to be consistent with docker image.	6 months ago
centralized_collection	run-pktvisor.sh IFS fix for args with spaces (#111)	2 months ago
cmake	feature/deps (#36)	7 months ago
cmd	require a new key 'kind' to specify policy type, in preparation for f...	13 days ago
docker	require binary arg to appimage to be consistent with docker image.	6 months ago
docs	rfcs	5 months ago
golang	feature/cli refactor (#120)	10 days ago
integration_tests	implement #78 dns filters (#105)	3 months ago
src	require a new key 'kind' to specify policy type, in preparation for f...	13 days ago
.clang-format	import	2 years ago
.dockerignore	issue #94 add TLS support to web server	4 months ago
.gitignore	Merge branch 'release' into feature/3.2.1-merge	4 months ago
.gitmodules	Modularize #23 (#27)	7 months ago
CMakeLists.txt	merge 3.2.0 release, go 3.3.0-develop	5 months ago
CODE_OF_CONDUCT.md	Create CODE_OF_CONDUCT.md	last month
CONTRIBUTING.md	Improve READMEs, other minor improvements (#25)	9 months ago

About

pktvisor is an observability agent that summarizes deep network data streams in real time, enabling on-node and centralized data visibility and analysis

[pktvisor.dev](#)

agent monitoring grafana prometheus observability packet-capture api-first data-streams collector-agent datasketches stream-processors stream-summarization

Readme

MPL-2.0 License

Releases 10

3.2.1 Latest on Jun 13

+ 9 releases

Contributors 8

pktvisor+Prometheus



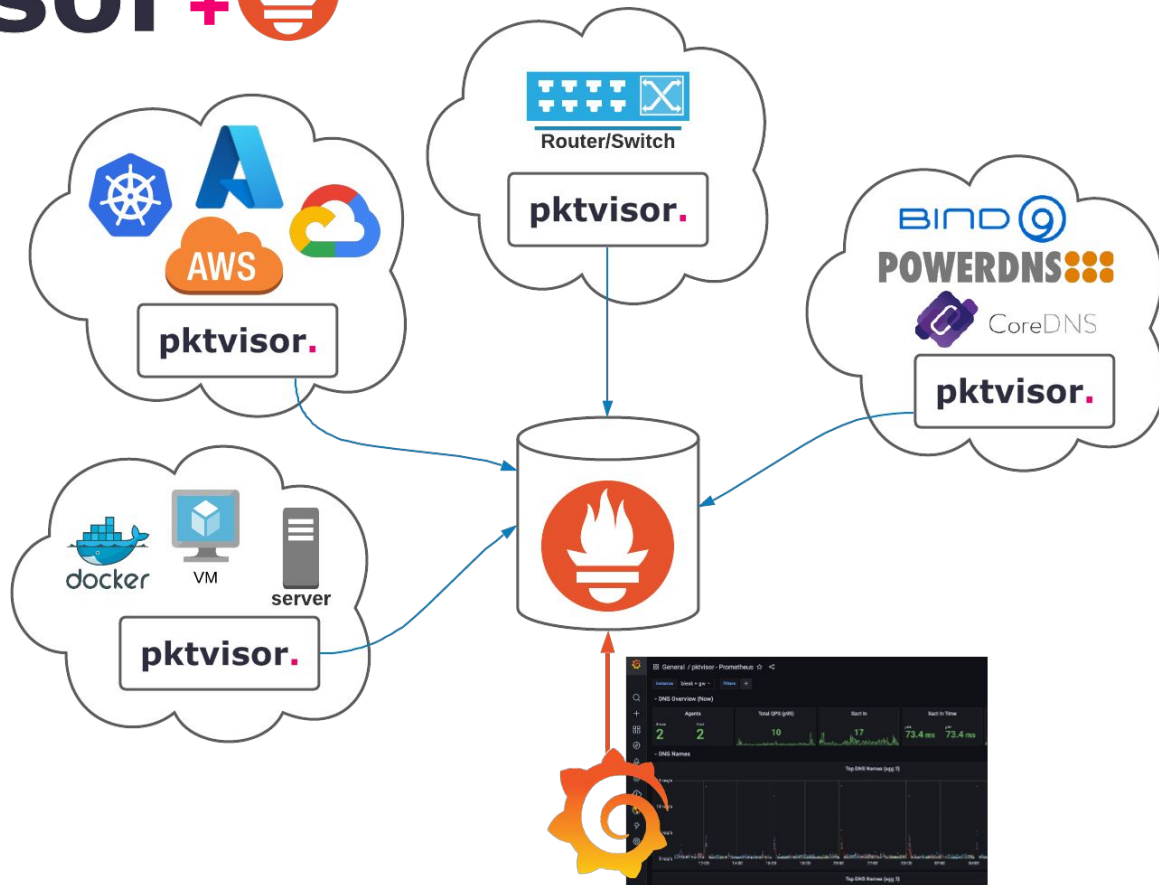
The Big Picture: pktvisor+Prometheus



PromCon
North America 2021

- ④ pktvisor exposes **Prometheus** metrics **per policy**
- ④ Provides a **global view** of distributed pktvisor agents
- ④ May be **scraped** or push with **remote write**
- ④ Grafana and other tools for exploring, **visualizing** and **alerting**

pktvisor+



Easy Remote Write



PromCon
North America 2021

pktvisor



Prometheus



Grafana Agent



ns1labs/pktvisor-prom-write ☆

↓ Pulls 436

By [ns1labs](#) • Updated 17 hours ago

pktvisor with native ability to send metrics to Prometheus through remote write

Container

Overview

Tags

pktvisor + centralized Prometheus collection

This container combines pktvisord with the [Grafana Agent](#) for collecting and sending metrics to Prometheus through remote write, including to cloud providers like [Grafana Cloud](#).

There is a sample [Grafana dashboard](#) which provides a good starting point for visualizing pktvisor metrics. You can also find it online via the [Grafana community dashboards](#), allowing you to import easily into any Grafana installation (ID 14221).

Example:

```
docker pull ns1labs/pktvisor-prom-write
docker run -d --net=host --env PKTVISORD_ARGS="--prom-instance <INSTANCE> <INTERFACE>" \
--env REMOTE_URL="https://<REMOTEHOST>/api/prom/push" --env USERNAME="<USERNAME>" \
--env PASSWORD="<PASSWORD>" ns1labs/pktvisor-prom-write
```

Grafana Dashboard



PromCon
North America 2021

[Grafana](#)[Products](#)[Open Source](#)[Learn](#)[Downloads](#)[Contact us](#)[Login](#)

All dashboards » [pktvisor - prometheus](#)



pktvisor - prometheus by [ns1labs](#)

 DASHBOARD

A dashboard for pktvisor observability tool (<https://github.com/ns1labs/pktvisor>), showcasing Network and DNS metrics.

Last updated: a month ago

Downloads:

Reviews:

Add your review!

Start with Grafana Cloud and the new FREE tier. Includes 10K series Prometheus or Graphite Metrics and 50gb Loki Logs

Overview

Revisions

Reviews




[pktvisor](#) summarizes network data streams in real time. It can capture Network, DNS, and other metrics via packet capture, dnstap, sflow, and other input methods.

This dashboard can be used as a starting point to visualize pktvisor metrics. See the Github page for information on how to deploy and collect these metrics.

Get this dashboard:


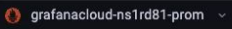
14221

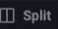

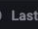

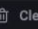

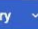
 Copy ID to Clipboard


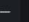
[Download JSON](#)

[How do I import this dashboard?](#)

Dependencies:

 Explore 

 Split   Last 1 hour   Clear all  Run query 


Metrics browser > Enter a PromQL query (run with Shift+Enter) 0.2s  

Query type


Range

Instant

Both


Step 


auto

Exemplars 

Help >

+ Add query

 Query history

 Inspector



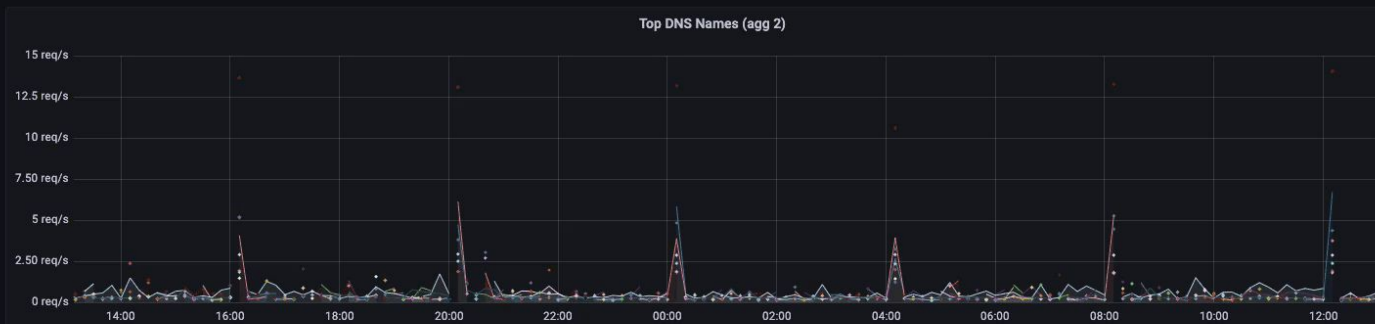
Instance blesk + gw ▾ Filters +

[pktvisor docs](#)

DNS Overview (Now)

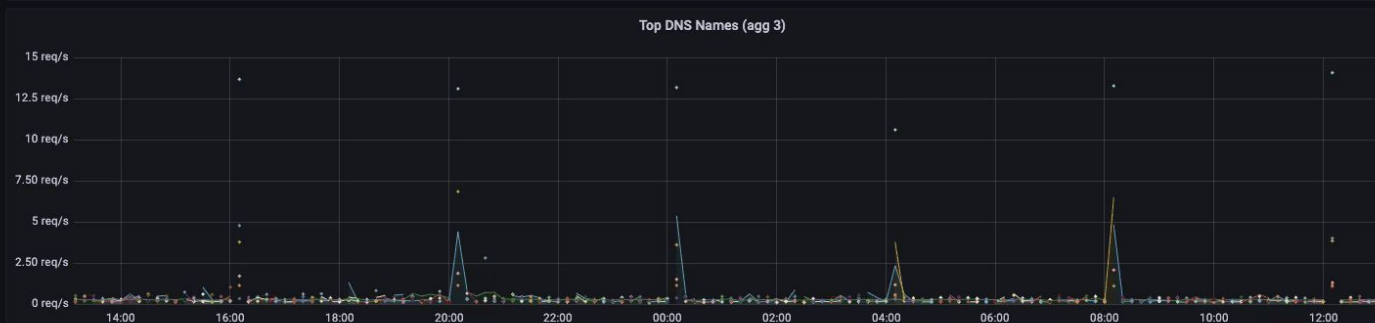
Agents		Total QPS (p95)	Xact In	Xact In Time		Xact Out	Xact Out Time		Unique QNames
Shown	Total			p95	p99		p95	p99	
2	2	19	27	71.6 ms	71.6 ms	50	85.3 ms	85.3 ms	45

DNS Names



Names (agg 2)

QName 🔍	Requests (sum) + 🔍
.google.com	20.9 K
.roku.com	19.5 K
.googleapis.com	7.23 K
.grafana.net	6.95 K
.akadns.net	6.43 K
.apple.com	5.22 K
.akamaiedge.net	5.18 K



Names (agg 3)

QName 🔍	Requests (sum) + 🔍
.logs.roku.com	18.7 K
.com.akadns.net	6.07 K
ns1rd81.grafana.net	5.42 K
.clients6.google.com	4.68 K
.oca.nflxvideo.net	4.67 K
.l.google.com	3.65 K
play.google.com	2.77 K

DNS

DNS Packets (In/Out)

QTime

Result Codes

Transport



The Big Picture: Orb



PromCon
North America 2021

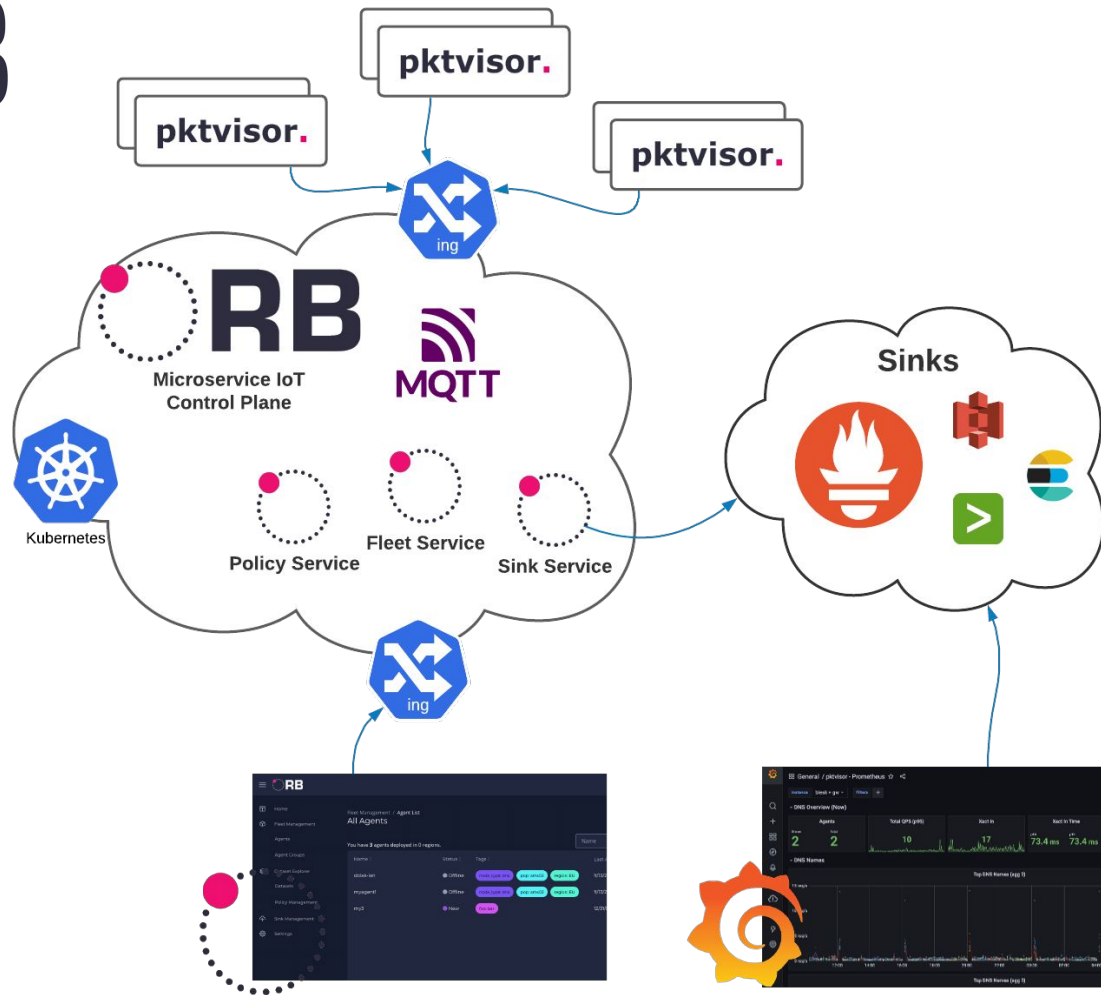
- 🔥 Nascent **free** and **open source** observability **platform** backed by NSI
- 🔥 Based on **IoT** principals, provides **UI, API** and **agent communication**
- 🔥 Solves challenges
 - **orchestrating** agents and their **policies**
 - **collecting** and **sinking** agent output

Orb Features



PromCon
North America 2021

- 🌀 **Multi-tenant fleet management** for agents
- 🌀 Dynamic **agent grouping** (based on tagging) for assigning policies
- 🌀 Definition and **orchestration of policies**, updating fleet in real time
- 🌀 Centralized **metric collection** and **sinking** to **multiple destinations** per policy
- 🌀 Helm chart for self-host **k8s deploy, free SaaS** coming soon at <http://orb.live>
- 🌀 **Modular** support for observability agents (contributors!)



- Home
- Fleet Management
- Agents
- Agent Groups
- Dataset Explorer
- Datasets
- Policy Management
- Sink Management
- Settings

Fleet Management / Agent List

All Agents

You have **3** agents deployed in 0 regions.

Name



Search by name

Name

Status

Tags

Last Activity

[+ NEW AGENT](#)

oblak-lan

● Offline

node_type: dns

pop: ams02

region: EU

9/13/21, 1:32 PM



myagent1

● Offline

node_type: dns

pop: ams02

region: EU

9/13/21, 1:38 PM



my3

● New

foo: bar

12/31/01, 7:03 PM





- FOSS
- Contributors welcome
- Please star!

ns1labs / orb

Public

Unwatch

8

Unstar

39

Fork

3

<> Code

Issues 53

Pull requests 6

Discussions

Actions

Projects 3

Wiki

Security

Insights

Settings

develop

20 branches

0 tags

Go to file

Add file

Code

dscabral

Merge pull request #335 from ns1labs/feature/ORB-edit-dataset

42c81d0

1 hour ago

822 commits

.github	feature/sinker (#307)	6 days ago
.idea/runConfigurations	hotfix(orb_policies): add ide fleet config for main flux sdk	4 days ago
RFCs	separate policies from datasets, add to repo. and policy repo tests.	3 months ago
agent	feature/sinker (#307)	6 days ago
cmd	feature/sinker (#307)	6 days ago
docker	feature/sinker (#307)	6 days ago
docs/images	update arch diagrams	6 days ago
fleet	Merge pull request #335 from ns1labs/feature/ORB-edit-dataset	1 hour ago
grafana/api	Postman updated orb-grafana 1.0.0	3 months ago
internal/httputil	feat(agent_group): add feature to filter agent groups by tags	last month
pkg	feature/sinker (#307)	6 days ago
policies	Merge pull request #335 from ns1labs/feature/ORB-edit-dataset	1 hour ago
sinker	feature/sinker (#307)	6 days ago
sinks	feature/sinker (#307)	6 days ago
ui	Merge branch 'develop' into fix/orb-ui-update-cache-upon-deleting	3 days ago
.dockerignore	chore(boilerplate): added mainflux/ui boilerplate to Orb	3 months ago
.gitignore	feat(mock-server): adjusted settings	2 months ago
CODE_OF_CONDUCT.md	Create CODE_OF_CONDUCT.md	last month

About

Orb is a cloud native orchestration platform for dynamic edge observability

[getorb.io](#)

kubernetes

iot

ui

docker-compose

metrics

self-hosted

cloud-native

control-plane

observability

fleet-management

edge-computing

Readme

MPL-2.0 License

Releases

No releases published

Create a new release

Contributors 11

Languages

Summary

The Big Picture



PromCon
North America 2021

pktvisor and **Orb**
supplement modern
observability stacks
by facilitating
**edge network
observability**

The projects are **free**
and **open source**
backed by **NSI**.

The ultimate
goal is **dynamic
orchestration** of
business intelligence
at the **edge**

Thank You!



PromCon

North America 2021

<https://pktvisor.dev/>

<https://getorb.io/>

🔗 sweyrick@ns1.com

🔗 [Orb Announcement List](#)

🔗 [NS1 Labs Slack](#)

🔗 In person demos at NS1 booth Wed - Fri

🔗 Or contact me on Virtual Platform
or CNCF Slack!

